

Pipeline Internals

Pipeline Internals I

This supplemental section documents the data structures and on-disk artifacts the pipeline produces, for readers who want to extend or audit it.

Data structures I

The Mermaid class diagram in this subsection shows the canonical fields each record carries through the pipeline. Records have additional optional metadata (e.g. `Paper.url`, `Paper.publisher`, `Paper.isbn`, `Paper.raw`) omitted for readability — consult `infra/structure/search/literature/models.py` (source on GitHub) and `src/pipeline.py` (source on GitHub) for the full schema.

```
classDiagram
    class SearchQuery {
        +str text
        +int max_results
        +int|None year_min
        +int|None year_max
        +list~str~ sources
        +list~str~ fields
        +matches_year(year) bool
    }
```

Data structures II

```
class Paper {  
    +str id  
    +str title  
    +list~str~ authors  
    +int|None year  
    +str|None doi  
    +str|None url  
    +str|None abstract  
    +str|None pdf_url  
    +str|None fulltext  
    +str|None venue  
    +str|None venue_type  
    +str|None volume  
    +str|None issue  
    +str|None pages  
    +str|None publisher
```

Data structures III

```
+str|None edition
+str|None isbn
+list~str~ keywords
+str|None source
+float score
+dict raw
+to_dict() dict
+from_dict(d) Paper
}

class SearchResult {
    +SearchQuery query
    +list~Paper~ papers
    +dict per_source_counts
    +dict errors
    +to_dict() dict
    +to_json() str
```

Data structures IV

```
}
```

```
class BibEntry {  
    +str entry_type  
    +str citation_key  
    +OrderedDict fields  
}
```

```
class BibDatabase {  
    +list~BibEntry~ entries  
    +str|None preamble  
    +find(key) BibEntry  
}
```

```
class LiteratureRunArtifacts {  
    +SearchResult result  
    +Path|None corpus_path
```

Data structures V

```
+Path|None bibtex_path
+list~FetchResult~ enrichment_log
+Path|None cache_dir
+dict citation_keys
+papers() list~Paper~
}
```

```
class ManuscriptVariables {
    +str config_query
    +int config_max_results
    +str config_sources
    +int result_num_papers
    +int result_num_sources
    +str result_per_source
    +str result_errors
    +str result_year_min
    +str result_year_max
}
```

Data structures VI

```
+int result_with_abstract
+int result_with_doi
+int deep_max_results_per_keyword
+int deep_keyword_count
+str deep_keywords_joined
+str deep_sources
+str deep_unique_papers
+as_dict() dict
+as_uppercase_keys() dict
}

class SynthesisResult {
    +str kind
    +str prompt
    +str text
    +str|None paper_id
}
```


Data structures VII

SearchResult --> SearchQuery

SearchResult --> "*" Paper

LiteratureRunArtifacts --> SearchResult

LiteratureRunArtifacts --> "*" BibEntry : via paper_to_

BibDatabase --> "*" BibEntry

On-disk layout I

The project keeps committed input data in `data/`, regeneratable outputs in `output/` (gitignored), and the manuscript source in `manuscript/`. The Mermaid flowchart in this subsection (rendered in the HTML build; the PDF build strips Mermaid) lists every artefact the standard pipeline writes:

- ▶ `data/corpus.json` — the bundled offline corpus, CI-safe default for LocalBackend.
- ▶ `output/search/results.json` — raw `SearchResult` JSON from the latest run.
- ▶ `output/search/cache/search_<hash>.json` — deterministic `SearchCache` files.
- ▶ `output/cache/abs/<safe_id>.txt` — one cached abstract per paper.
- ▶ `output/cache/pdf/<safe_id>.{pdf,txt}` — PDF bytes plus extracted text.

On-disk layout II

- ▶ `output/llm/synthesis.md` — corpus-level LLM synthesis (when enabled).
- ▶ `output/llm/per_paper/<safe_id>.md` — one per-paper note per paper (when enabled).
- ▶ `output/figures/{papers_per_source,year_histogram,score_`
— diagnostic figures.
- ▶ `output/data/manuscript_variables.json` — substitution table consumed by the resolver.
- ▶ `output/corpus.json` — enriched corpus, written in LocalBackend-compatible format so it can re-seed a deterministic future run.
- ▶ `output/enrichment_log.json` — one entry per fetcher per paper.
- ▶ `output/reading_report.md` — final markdown reading report.

On-disk layout III

- ▶ `output/run_summary.json` — one-line metadata for the run.
- ▶ `manuscript/references.bib` — auto-populated, Pandoc-ready BibTeX.

flowchart TB

```
P["projects/templates/template_search_project/"]  
P --> DAT["data/<br/>committed"]  
P --> OUT["output/<br/>regeneratable · gitignored"]  
P --> MAN["manuscript/"]
```

```
DAT --> CORP_IN["corpus.json<br/>bundled offline corpus"]
```

```
OUT --> SR["search/"]  
OUT --> CC["cache/"]  
OUT --> LD["llm/"]  
OUT --> FIG["figures/"]  
OUT --> DD["data/"]
```

On-disk layout IV

```
OUT --> CORP_OUT[corpus.json<br/>enriched · LocalBacker
OUT --> ELOG[enrichment_log.json<br/>per-fetch status]
OUT --> RR[reading_report.md]
OUT --> RS[run_summary.json]

SR --> SR_R[results.json<br/>SearchResult]
SR --> SR_C["cache/search_HASH.json<br/>deterministic S

CC --> CC_A["abs/SAFE_ID.txt"]
CC --> CC_P["pdf/SAFE_ID.pdf and .txt"]

LD --> LD_S[synthesis.md]
LD --> LD_PP["per_paper/SAFE_ID.md"]

FIG --> F1[papers_per_source.png]
FIG --> F2[year_histogram.png]
FIG --> F3[score_distribution.png]
```

On-disk layout V

```
DD --> MV[manuscript_variables.json]
```

```
MAN --> BIB[references.bib<br/>auto-populated · Pandoc-
```

```
classDef dir fill:#0f172a,stroke:#0f172a,color:#fff
```

```
classDef gen fill:#0f766e,stroke:#0f172a,color:#fff
```

```
classDef src fill:#1e3a8a,stroke:#0f172a,color:#fff
```

```
class P,DAT,OUT,MAN,SR,CC,LD,FIG,DD dir
```

```
class CORP_OUT,ELOG,RR,RS,SR_R,SR_C,CC_A,CC_P,LD_S,LD_F
```

```
class CORP_IN src
```

Citation-key collision handling I

`paper_to_bibentry()` generates citation keys as `<author><year><title-word>` (with stop-words filtered and unicode folded). When two papers in the same result set produce the same key — common when one author publishes multiple papers in the same year on closely related topics — `src/pipeline.py::_disambiguate_citation_key` appends a deterministic suffix from the alphabet (a, b, ..., z, then two-letter combinations aa, ab, ...) until uniqueness is restored, with a numeric `_1`, `_2`, ... fallback for the pathological case. The mapping is exposed to downstream stages via `LiteratureRunArtifacts.citation_keys`, and the report uses these keys verbatim, so the LLM synthesis and the BibTeX file always agree.

The deep-search workflow has its own collision handler in `src/deep_search.py::run_deep_search` that operates over the post-deduplication aggregate roster — see `[@sec:deep_search]` — and the unified `references_deep.bib` reflects the same mapping.

Failure isolation I

- ▶ **A backend is unreachable.** LiteratureClient records the message into `result.errors[name]` and continues. The reading report surfaces these errors in a callout block at the top.
- ▶ **An abstract fetch fails.** The fetcher records `status="error"` in `enrichment_log.json` and the paper keeps its existing (possibly empty) abstract.
- ▶ **A PDF fetch fails or pypdf is missing.** Same pattern — the PDF, if downloaded, is still cached on disk. The reading report does not reference the missing fulltext.
- ▶ **The LLM is unreachable or unconfigured.** `scripts/run_search_pipeline.py` logs a warning and skips the synthesis stage entirely — `output/llm/` is left empty and the reading report omits both the per-paper-notes and cross-corpus sections. No placeholder text is ever written into the archive, so a missing LLM is observable from the absence of those sections rather than from a fake “(LLM unavailable)” string.

Failure isolation II