

Deep Search

Deep Search I

The deep-search workflow extends the standard literature pipeline along three axes: **breadth** (multi-keyword fan-out), **depth** (full enrichment of every paper, abstract + PDF fulltext), and **archival quality** (per-paper multi-section LLM reading notes saved as standalone markdown). It is invoked from the same configuration file but a separate orchestrator script.

Configuration I

The `deep_search:` block in `manuscript/config.yaml` controls the run. The most-used knobs:

- ▶ `keywords` — list of free-text queries (each becomes one `SearchQuery`).
- ▶ `max_results_per_keyword` — per-keyword cap (100 by default; honoured by the aggregator after dedup).
- ▶ `sources` — backend list, same vocabulary as the standard pipeline (`arxiv`, `crossref`, `local`, `paperclip`).
- ▶ `fetch_abstracts` / `fetch_fulltext` — when both are true, every returned paper has its abstract fetched (arXiv export API) and (where a PDF URL is available) its fulltext extracted via `pypdf`. Cached on disk under `output/cache/abstracts/` and `output/cache/pdf/`.

Configuration II

- ▶ `llm_per_paper` — when true and Ollama is reachable, each paper gets a multi-section markdown reading note generated by the local LLM. When false, or when the LLM stack is genuinely unreachable at runtime, the per-paper note is written with only the abstract / fulltext-excerpt sections; the synthesis section is omitted entirely (no placeholder text). The composer's Supplemental S1 also drops the per-paper synthesis subsection in that case rather than emitting empty rows.
- ▶ `write_unified_bibtex / unified_bibtex_path` — when true, a deduplicated, collision-suffixed BibTeX file is written under `manuscript/references_deep.bib` so it can be cited from the manuscript via Pandoc `[@key]` syntax exactly the same way as the hand-curated `references.bib`.

Pipeline shape I

The deep-search pipeline reads the `deep_search:` block from `manuscript/config.yaml`, runs one `SearchQuery` per keyword (capped at `max_results_per_keyword`), aggregates the per-backend results through `LiteratureClient`, enriches every paper via `AbstractFetcher` and `FulltextFetcher`, generates collision-free citation keys with `paper_to_bibentry`, and (when `lm_per_paper: true` and Ollama is reachable) calls the local LLM with `DEEP_PROMPT` to produce a seven-section reading note per paper. The per-keyword outputs are then merged by `merge_papers` to form a deduplicated aggregate roster which is written to `manuscript/references_deep.bib`, `output/deep_search/aggregate.json`, and `output/deep_search/aggregate_report.md`. A Mermaid flowchart in this subsection renders the same flow visually in the HTML build.

Pipeline shape II

flowchart TB

CFG[deep_search block of config.yaml]

CFG --> KW{{For each keyword}}

KW --> SEARCH[SearchQuery · max_results_per_keyword]

SEARCH --> CLIENT[LiteratureClient
aggregates arxiv]

CLIENT --> ENRICH[AbstractFetcher + FulltextFetcher
]

ENRICH --> KEYS[paper_to_bibentry
collision-free ke

KEYS --> LLMQ{LLM per-paper?}

LLMQ -- yes --> LLM[Ollama LLMClient
DEEP_PROMPT ·

LLMQ -- no --> SKIP[skip synthesis · no placeholder]

LLM --> NOTE[output/deep_search/<slug>/per_paper/

SKIP --> NOTE

NOTE --> KR[KeywordResult]

KR --> AGG[merge_papers across all keywords]

AGG --> BIB[unified BibTeX
references_deep.bib]

AGG --> AJSON[aggregate.json]

Pipeline shape III

AGG --> AMD[aggregate_report.md]

```
classDef cfg fill:#0f172a,stroke:#0f172a,color:#fff
classDef proc fill:#1e3a8a,stroke:#0f172a,color:#fff
classDef llm fill:#7c2d12,stroke:#0f172a,color:#fff
classDef out fill:#0f766e,stroke:#0f172a,color:#fff
class CFG,KW,LLMQ cfg
class SEARCH,CLIENT,ENRICH,KEYS,KR,AGG proc
class LLM,SKIP llm
class NOTE,BIB,AJSON,AMD out
```

LLM prompt I

The deep-search prompt is richer than the standard synthesis.PROMPT_PER_PAPER. It produces a self-contained reading note suitable for archival without re-reading the paper:

Contribution

One paragraph stating the paper's central claim and why it

Method

2-4 bullets describing the technical approach.

Evidence

2-4 bullets describing experiments / proofs.

Limitations

1-3 bullets covering caveats and what the paper does NOT do

Connections

LLM prompt II

1-3 bullets relating this paper to other work in the field
papers explicitly named in the input.

Significance for {keyword}

One paragraph explaining why this paper matters for the key

Tags

5-10 lowercase keywords.

The LLM is duck-typed as `Callable[[str], str]` so tests pass a deterministic local callable that returns real, well-formed reading-note text; runtime callers wrap an `infrastructure.llm.LLMClient` with `seed=42`, `temperature=0.0` for reproducibility. When the LLM stack is unreachable, callers pass `None` and the per-paper synthesis stage is skipped entirely — no placeholder text is ever written into the archive.

On-disk layout I

A successful deep-search run writes the following artefacts (paths shown relative to the project root); a Mermaid tree in this subsection renders the same hierarchy in the HTML build:

- ▶ `output/deep_search/aggregate.json` — every unique paper across keywords.
- ▶ `output/deep_search/aggregate_report.md` — cross-keyword markdown summary.
- ▶ `output/deep_search/run_summary.json` — one-line metadata for the run.
- ▶ `output/deep_search/<keyword_slug>/papers.json` — enriched per-keyword paper list.
- ▶ `output/deep_search/<keyword_slug>/reading_report.md` — per-keyword markdown summary.

On-disk layout II



output/deep_search/<keyword_slug>/per_paper/<safe_id>.m
— one reading note per paper.



manuscript/references_deep.bib — auto-generated,
deduplicated unified BibTeX file.

flowchart TB

```
P["projects/templates/template_search_project/"]
```

```
P --> OUT["output/deep_search/"]
```

```
P --> BIB[manuscript/references_deep.bib<br/>auto-generated]
```

```
OUT --> AGG[aggregate.json<br/>every unique paper]
```

```
OUT --> AMD[aggregate_report.md<br/>cross-keyword summary]
```

```
OUT --> RS[run_summary.json<br/>one-line metadata]
```

```
OUT --> KW1["<keyword_slug_1>"]
```

```
OUT --> KW2["<keyword_slug_2>"]
```

```
OUT --> KW3["<keyword_slug_3>"]
```

On-disk layout III

```
KW1 --> KW1_P[papers.json]
KW1 --> KW1_R[reading_report.md]
KW1 --> KW1_PP["per_paper/<br/>SAFE_ID.md per paper"]
```

```
classDef d fill:#0f172a,stroke:#0f172a,color:#fff
classDef f fill:#0f766e,stroke:#0f172a,color:#fff
classDef bib fill:#1e3a8a,stroke:#0f172a,color:#fff
class P,OUT,KW1,KW2,KW3,KW1_PP d
class AGG,AMD,RS,KW1_P,KW1_R f
class BIB bib
```

Determinism I

Three caches make this workflow replayable:

1. SearchCache
(output/search/cache/search_<hash>.json) — keyed on canonical query identity per keyword.
2. Abstract cache (output/cache/abs/<safe_id>.txt).
3. Fulltext cache (output/cache/pdf/<safe_id>.{pdf,txt}).
4. LLM seed (deep_search.llm_seed: 42) + temperature (0.0).

Commit any subset of these caches to version control to freeze a run.

Paperclip backend I

The Paperclip backend is opt-in. Enable it by:

1. Creating projects/templates/template_search_project/.env (gitignored) with your PAPERCLIP_API_KEY=gxl_... (template at .env.example).
2. Including paperclip in deep_search.sources.

The orchestrator scripts auto-load .env via the lightweight src.dotenv module before constructing the backend list. The PaperclipBackend adapter mirrors the wire protocol of the official gxl_paperclip Python SDK: POST /papers with an X-API-Key header and a JSON-RPC tools/call envelope. We support both the modern structuredContent.papers response shape and the older text-only content[].text shape (best-effort regex extraction).

Paperclip backend II

Failure-isolation: any per-backend error (including the migration-state HTTP 405 the production endpoint may currently return) is captured into `SearchResult.errors[paperclip]` and the run continues. See `output/deep_search/<keyword>/papers.json` and the aggregate `run_summary.json` for the recorded errors.

CLI I

Run with everything from config.yaml (assumes deep_search)

```
uv run python projects/templates/template_search_project/sc
```

Force-enable without touching config.yaml

```
uv run python projects/templates/template_search_project/sc
```

Override keyword list at the command line

```
uv run python projects/templates/template_search_project/sc  
    --enable --keyword "convex optimization" --keyword "sto
```

Skip LLM stage even when config enables it

```
uv run python projects/templates/template_search_project/sc
```

Bypass cache (writes still happen)

```
uv run python projects/templates/template_search_project/sc
```

Local-corpus mode (offline · CI-friendly)

CLI II

```
uv run python projects/templates/template_search_project/s  
    --enable --corpus projects/templates/template_search_pr
```